

PICT, PUNE  
SL-VI Assignment 4

Version-Hadoop-2.7.2

Version-Hive-2.2.0

Version-Hbase-1.2.6

4. Write an application using HBase and HiveQL for flight information system which will include

- a) Creating, Dropping, and altering Database tables
- b) Creating an external Hive table to connect to the HBase for Customer Information Table
- c) Load table with data, insert new values and field in the table, Join tables with Hive
- d) Create index on Flight information Table
- e) Find the average departure delay per day in 2008.

### **a) Creating, Dropping, and altering Database tables Using Hbase**

#### **#Create Table:**

```
hbase(main):002:0> create 'flight','finfo','fsch'
```

```
0 row(s) in 4.6960 seconds
```

```
=> Hbase::Table - flight
```

#### **#Table Created-list**

```
hbase(main):003:0> list
```

```
TABLE
```

```
flight
```

```
table1
```

```
table2
```

```
3 row(s) in 0.0120 seconds
```

#### **#Insert records in created table**

```
hbase(main):004:0> put 'flight',1,'finfo:source','pune'
```

```
0 row(s) in 0.2480 seconds
```

```
hbase(main):008:0> put 'flight',1,'finfo:dest','mumbai'
```

```
0 row(s) in 0.0110 seconds
```

```
hbase(main):010:0> put 'flight',1,'fsch:at','10.25a.m.'
```

```
0 row(s) in 0.0060 seconds
```

```
hbase(main):011:0> put 'flight',1,'fsch:dt','11.25 a.m.'
0 row(s) in 0.0070 seconds
hbase(main):012:0> put 'flight',1,'fsch:delay','5min'
hbase(main):015:0> put 'flight',2,'finfo:source','pune'
0 row(s) in 0.0160 seconds
hbase(main):016:0> put 'flight',2,'finfo:dest','kolkata'
0 row(s) in 0.0070 seconds
hbase(main):017:0> put 'flight',2,'fsch:at','7.00a.m.'
0 row(s) in 0.0080 seconds
```

```
hbase(main):018:0> put 'flight',2,'fsch:dt','7.30a.m.'
0 row(s) in 0.0050 seconds
hbase(main):019:0> put 'flight',2,'fsch:delay','2 min'
0 row(s) in 0.0090 seconds
hbase(main):021:0> put 'flight',3,'finfo:source','mumbai'
0 row(s) in 0.0040 seconds
hbase(main):022:0> put 'flight',3,'finfo:dest','pune'
0 row(s) in 0.0070 seconds
hbase(main):023:0> put 'flight',3,'fsch:at','12.30p.m.'
0 row(s) in 0.0100 seconds
hbase(main):024:0> put 'flight',3,'fsch:dt','12.45p.m.'
0 row(s) in 0.0040 seconds
hbase(main):025:0> put 'flight',3,'fsch:delay','1 min'
0 row(s) in 0.0190 seconds
hbase(main):026:0> put 'flight',4,'finfo:source','mumbai'
0 row(s) in 0.0060 seconds
hbase(main):027:0> put 'flight',4,'finfo:dest','delhi'
0 row(s) in 0.0050 seconds
hbase(main):028:0> put 'flight',4,'fsch:at','2.00p.m.'
0 row(s) in 0.0080 seconds
hbase(main):029:0> put 'flight',4,'fsch:dt','2.45p.m.'
0 row(s) in 0.0040 seconds
hbase(main):030:0> put 'flight',4,'fsch:delay','10 min'
0 row(s) in 0.0140 seconds
```

#### **#Display Records from Table 'flight'**

```
hbase(main):031:0> scan 'flight'
ROW          COLUMN+CELL
1            column=finfo:dest, timestamp=1521312730758, value=mumbai
1            column=finfo:source, timestamp=1521312493881, value=pune
1            column=fsch:at, timestamp=1521312789417, value=10.25a.m.
1            column=fsch:delay, timestamp=1521312850594, value=5min
```

```

1      column=fsch:dt, timestamp=1521312823256, value=11.25 a.m.
2      column=finfo:dest, timestamp=1521313135697, value=kolkata
2      column=finfo:source, timestamp=1521313092772, value=pune
2      column=fsch:at, timestamp=1521313166540, value=7.00a.m.
2      column=fsch:delay, timestamp=1521313229963, value=2 min
2      column=fsch:dt, timestamp=1521313202767, value=7.30a.m.
3      column=finfo:dest, timestamp=1521313310302, value=pune
3      column=finfo:source, timestamp=1521313290906, value=mumbai
3      column=fsch:at, timestamp=1521313333432, value=12.30p.m.
3      column=fsch:delay, timestamp=1521313379725, value=1 min
3      column=fsch:dt, timestamp=1521313353804, value=12.45p.m.
4      column=finfo:dest, timestamp=1521313419679, value=delhi
4      column=finfo:source, timestamp=1521313404831, value=mumbai
4      column=fsch:at, timestamp=1521313440328, value=2.00p.m.
4      column=fsch:delay, timestamp=1521313472389, value=10 min
4      column=fsch:dt, timestamp=1521313455226, value=2.45p.m.
4 row(s) in 0.0300 seconds

```

#### #Alter Table (add one more column family)

```
hbase(main):036:0> alter 'flight',NAME=>'revenue'
```

Updating all regions with the new schema...

0/1 regions updated.

1/1 regions updated.

Done.

0 row(s) in 3.7640 seconds

```
hbase(main):037:0> scan 'flight'
```

ROW	COLUMN+CELL
1	column=finfo:dest, timestamp=1521312730758, value=mumbai
1	column=finfo:source, timestamp=1521312493881, value=pune
1	column=fsch:at, timestamp=1521312789417, value=10.25a.m.
1	column=fsch:delay, timestamp=1521312850594, value=5min
1	column=fsch:dt, timestamp=1521312823256, value=11.25 a.m.
2	column=finfo:dest, timestamp=1521313135697, value=kolkata
2	column=finfo:source, timestamp=1521313092772, value=pune
2	column=fsch:at, timestamp=1521313166540, value=7.00a.m.
2	column=fsch:delay, timestamp=1521313229963, value=2 min
2	column=fsch:dt, timestamp=1521313202767, value=7.30a.m.
3	column=finfo:dest, timestamp=1521313310302, value=pune
3	column=finfo:source, timestamp=1521313290906, value=mumbai
3	column=fsch:at, timestamp=1521313333432, value=12.30p.m.
3	column=fsch:delay, timestamp=1521313379725, value=1 min
3	column=fsch:dt, timestamp=1521313353804, value=12.45p.m.
4	column=finfo:dest, timestamp=1521313419679, value=delhi
4	column=finfo:source, timestamp=1521313404831, value=mumbai
4	column=fsch:at, timestamp=1521313440328, value=2.00p.m.
4	column=fsch:delay, timestamp=1521313472389, value=10 min

4 column=fsch:dt, timestamp=1521313455226, value=2.45p.m.  
4 row(s) in 0.0290 seconds

#### #Insert records into added column family

hbase(main):038:0> put 'flight',4,'revenue:rs','45000'  
0 row(s) in 0.0100 seconds

#### #Check the updates

hbase(main):039:0> scan 'flight'

ROW	COLUMN+CELL
1	column=finfo:dest, timestamp=1521312730758, value=mumbai
1	column=finfo:source, timestamp=1521312493881, value=pune
1	column=fsch:at, timestamp=1521312789417, value=10.25a.m.
1	column=fsch:delay, timestamp=1521312850594, value=5min
1	column=fsch:dt, timestamp=1521312823256, value=11.25 a.m.
2	column=finfo:dest, timestamp=1521313135697, value=kolkata
2	column=finfo:source, timestamp=1521313092772, value=pune
2	column=fsch:at, timestamp=1521313166540, value=7.00a.m.
2	column=fsch:delay, timestamp=1521313229963, value=2 min
2	column=fsch:dt, timestamp=1521313202767, value=7.30a.m.
3	column=finfo:dest, timestamp=1521313310302, value=pune
3	column=finfo:source, timestamp=1521313290906, value=mumbai
3	column=fsch:at, timestamp=1521313333432, value=12.30p.m.
3	column=fsch:delay, timestamp=1521313379725, value=1 min
3	column=fsch:dt, timestamp=1521313353804, value=12.45p.m.
4	column=finfo:dest, timestamp=1521313419679, value=delhi
4	column=finfo:source, timestamp=1521313404831, value=mumbai
4	column=fsch:at, timestamp=1521313440328, value=2.00p.m.
4	column=fsch:delay, timestamp=1521313472389, value=10 min
4	column=fsch:dt, timestamp=1521313455226, value=2.45p.m.
4	column=revenue:rs, timestamp=1521314406914, value=45000

4 row(s) in 0.0340 seconds

#### #Delete Column family

hbase(main):040:0> alter 'flight',NAME=>'revenue',METHOD=>'delete'  
Updating all regions with the new schema...  
0/1 regions updated.  
1/1 regions updated.  
Done.  
0 row(s) in 3.7880 seconds

#### #changes Reflected in Table

hbase(main):041:0> scan 'flight'

ROW	COLUMN+CELL
1	column=finfo:dest, timestamp=1521312730758, value=mumbai
1	column=finfo:source, timestamp=1521312493881, value=pune
1	column=fsch:at, timestamp=1521312789417, value=10.25a.m.
1	column=fsch:delay, timestamp=1521312850594, value=5min
1	column=fsch:dt, timestamp=1521312823256, value=11.25 a.m.
2	column=finfo:dest, timestamp=1521313135697, value=kolkata

```

2      column=finfo:source, timestamp=1521313092772, value=pune
2      column=fsch:at, timestamp=1521313166540, value=7.00a.m.
2      column=fsch:delay, timestamp=1521313229963, value=2 min
2      column=fsch:dt, timestamp=1521313202767, value=7.30a.m.
3      column=finfo:dest, timestamp=1521313310302, value=pune
3      column=finfo:source, timestamp=1521313290906, value=mumbai
3      column=fsch:at, timestamp=1521313333432, value=12.30p.m.
3      column=fsch:delay, timestamp=1521313379725, value=1 min
3      column=fsch:dt, timestamp=1521313353804, value=12.45p.m.
4      column=finfo:dest, timestamp=1521313419679, value=delhi
4      column=finfo:source, timestamp=1521313404831, value=mumbai
4      column=fsch:at, timestamp=1521313440328, value=2.00p.m.
4      column=fsch:delay, timestamp=1521313472389, value=10 min
4      column=fsch:dt, timestamp=1521313455226, value=2.45p.m.
4 row(s) in 0.0280 seconds

```

## #Drop Table

### #Create Table for dropping

```

hbase(main):046:0> create 'tb1','cf'
0 row(s) in 2.3120 seconds
=> Hbase::Table - tb1

```

```

hbase(main):047:0> list
TABLE
flight
table1
table2
tb1
4 row(s) in 0.0070 seconds
=> ["flight", "table1", "table2", "tb1"]

```

## #Drop Table

```
hbase(main):048:0> drop 'tb1'
```

ERROR: Table tb1 is enabled. Disable it first.

Here is some help for this command:  
 Drop the named table. Table must first be disabled:  
 hbase> drop 't1'  
 hbase> drop 'ns1:t1'

## #Disable table

```

hbase(main):049:0> disable 'tb1'
0 row(s) in 4.3480 seconds

```

```

hbase(main):050:0> drop 'tb1'
0 row(s) in 2.3540 seconds

```

```

hbase(main):051:0> list
TABLE

```

flight  
table1  
table2  
3 row(s) in 0.0170 seconds

=> ["flight", "table1", "table2"]

#### **#Read data from table for row key 1:**

```
hbase(main):052:0> get 'flight',1
COLUMN          CELL
info:dest       timestamp=1521312730758, value=mumbai
info:source     timestamp=1521312493881, value=pune
fsch:at         timestamp=1521312789417, value=10.25a.m.
fsch:delay      timestamp=1521312850594, value=5min
fsch:dt         timestamp=1521312823256, value=11.25 a.m.
5 row(s) in 0.0450 seconds
```

#### **Read data for particular column from HBase table:**

```
hbase(main):053:0> get 'flight','1',COLUMN=>'info:source'
COLUMN          CELL
info:source     timestamp=1521312493881, value=pune
1 row(s) in 0.0110 seconds
```

#### **Read data for multiple columns in HBase Table:**

```
hbase(main):054:0> get 'flight','1',COLUMN=>['info:source','info:dest']
COLUMN          CELL
info:dest       timestamp=1521312730758, value=mumbai
info:source     timestamp=1521312493881, value=pune
2 row(s) in 0.0190 seconds
```

```
hbase(main):055:0> scan 'flight',COLUMNS=>'info:source'
ROW             COLUMN+CELL
1               column=info:source, timestamp=1521312493881, value=pune
2               column=info:source, timestamp=1521313092772, value=pune
3               column=info:source, timestamp=1521313290906, value=mumbai
4               column=info:source, timestamp=1521313404831, value=mumbai
4 row(s) in 0.0320 seconds
```

## **b) Creating an external Hive table to connect to the HBase for Customer Information Table**

**Covers==>**

### **c) Load table with data, insert new values and field in the table, Join tables with Hive**

### **Add these Jar files in Hive(on hive prompt)**

```
add jar file:///usr/local/HBase/lib/zookeeper-3.4.6.jar;
add jar file:///usr/local/HBase/lib/guava-12.0.1.jar;
add jar file:///usr/local/HBase/lib/hbase-client-1.2.6.jar;
add jar file:///usr/local/HBase/lib/hbase-common-1.2.6.jar;
add jar file:///usr/local/HBase/lib/hbase-protocol-1.2.6.jar;
add jar file:///usr/local/HBase/lib/hbase-server-1.2.6.jar;
add jar file:///usr/local/HBase/lib/hbase-shell-1.2.6.jar;
add jar file:///usr/local/HBase/lib/hbase-thrift-1.2.6.jar;
add jar file:///usr/local/hive/lib/hive-hbase-handler-2.2.0.jar;
```

### **Set the values of variables in Hive**

```
set hbase.zookeeper.quorum=localhost;
set hive.metastore.client.setugi=true;
set hive.exec.stagingdir=/tmp/.hivestage;
set hive.exec.dynamic.partition=true;
set hive.exec.dynamic.partition.mode=nonstrict;
set hive.auto.convert.join=false;
set hive.hbase.wal.enabled=false;
SET hive.exec.dynamic.partition = true;
SET hive.exec.dynamic.partition.mode = nonstrict;
SET hive.exec.max.dynamic.partitions = 10000;
SET hive.exec.max.dynamic.partitions.pernode = 1000;
```

### **# Create the external table emp using hive**

```
hive>create external table empdata2 ( ename string, esal int)
row format delimited fields terminated by "," stored as textfile location
"/home/hduser/Desktop/empdata2";
```

```
hive>load data local inpath '/home/hduser/Desktop/empdb.txt' into table empdata2;
```

### **#Create External Table in hive referring to hbase table**

#### **# create hbase table emphive first**

```
hbase(main):003:0> create 'emphive', 'cf'
0 row(s) in 4.6260 seconds
```

#### **#create hive external table**

```
CREATE external TABLE hive_table_emp(id int, name string, esal string)
STORED BY 'org.apache.hadoop.hive.hbase.HBaseStorageHandler'
WITH SERDEPROPERTIES ("hbase.columns.mapping" = ":key,cf:name,cf:esal")
TBLPROPERTIES ("hbase.table.name" = "emphive");
```

## # load data into hive\_table\_emp

(Hive doesn't allow directly inserting data into external hive table)

#for that create one hive table(managed table in hive)

**Managed table** and **External table in Hive**. There are two types of **tables in Hive**, one is **Managed table** and second is **external table**. the difference is, when you drop a **table**, if it is **managed table** **hive** deletes both data and meta data, if it is **external table** **Hive** only deletes metadata.

```
hive>create table empdbnew(eno int, ename string, esal int) row format delimited fields terminated by ',' stored as textfile;
```

## #load data in managed table

```
hive>load data local inpath '/home/hduser/Desktop/empdbnew.txt' into table empdbnew;
```

## #Load data in external table from managed table.

```
hive>INSERT INTO hive_table_emp select * from empdbnew;
```

```
hive> select * from hive_table_emp;
```

OK

```
1      deepali120000
2      mahesh      30000
3      mangesh     25000
4      ram         39000
5      brijesh     40000
6      john        300000
```

Time taken: 0.52 seconds, Fetched: 6 row(s)

## #display records where salary is greater than 40000

```
hive> select * from hive_table_emp where esal>40000;
```

OK

```
1      deepali120000
6      john        300000
```

Time taken: 0.546 seconds, Fetched: 2 row(s)

## #Check hbase for updates(The records are available in associated Hbase table)

```
hbase(main):008:0> scan 'emphive'
```

ROW	COLUMN+CELL
1	column=cf:esal, timestamp=1522212425665, value=120000
1	column=cf:name, timestamp=1522212425665, value=deepali
2	column=cf:esal, timestamp=1522212425665, value=30000
2	column=cf:name, timestamp=1522212425665, value=mahesh
3	column=cf:esal, timestamp=1522212425665, value=25000
3	column=cf:name, timestamp=1522212425665, value=mangesh
4	column=cf:esal, timestamp=1522212425665, value=39000
4	column=cf:name, timestamp=1522212425665, value=ram
5	column=cf:esal, timestamp=1522212425665, value=40000
5	column=cf:name, timestamp=1522212425665, value=brijesh
6	column=cf:esal, timestamp=1522212425665, value=300000
6	column=cf:name, timestamp=1522212425665, value=john



6 row(s) in 0.0700 seconds

## # Creating external table in Hive referring to Hbase

### #referring to flight table created in Hbase

```
CREATE external TABLE hbase_flight_new(fno int, fsource string,fdest string,fsh_at string,fsh_dt string,fsch_delay
string,delay int)
STORED BY 'org.apache.hadoop.hive.hbase.HBaseStorageHandler'
WITH SERDEPROPERTIES ("hbase.columns.mapping" =
":key,finfo:source,finfo:dest,fsch:at,fsch:dt,fsch:delay,delay:dl")
TBLPROPERTIES ("hbase.table.name" = "flight");
```

```
hive> CREATE external TABLE hbase_flight_new(fno int, fsource string,fdest string,fsh_at string,fsh_dt
string,fsch_delay string,delay int)
> STORED BY 'org.apache.hadoop.hive.hbase.HBaseStorageHandler'
> WITH SERDEPROPERTIES ("hbase.columns.mapping"
=":key,finfo:source,finfo:dest,fsch:at,fsch:dt,fsch:delay,delay:dl")
> TBLPROPERTIES ("hbase.table.name" = "flight");
OK
Time taken: 0.361 seconds
```

### #table created in hive

```
hive> show tables;
OK
abc
ddl_hive
emp
empdata
empdata1
empdata2
empdbnew
hbase_flight
hbase_flight1
hbase_flight_new
hbase_table_1
hive_table_emp
Time taken: 0.036 seconds, Fetched: 12 row(s)
```

### # Display records from that table

```
hive> select * from hbase_flight_new;
OK
1      pune   mumbai      10.25a.m.    11.25 a.m.    5min  10
2      pune   kolkata7.00a.m.    7.30a.m.    2 min  4
3      mumbai   pune   12.30p.m.    12.45p.m.    1 min  5
4      mumbai   delhi  2.00p.m.    2.45p.m.    10 min 16
Time taken: 0.581 seconds, Fetched: 4 row(s)
```

## e) Find the average departure delay per day in 2008.

### #calculate average delay

```
hive> select sum(delay) from hbase_flight_new;
```

WARNING: Hive-on-MR is deprecated in Hive 2 and may not be available in the future versions.  
Consider using a different execution engine (i.e. spark, tez) or using Hive 1.X releases.

Query ID = hduser\_20180328130004\_47384e9a-7490-4dfb-809d-ae240507bfab

Total jobs = 1

Launching Job 1 out of 1

Number of reduce tasks determined at compile time: 1

In order to change the average load for a reducer (in bytes):

set hive.exec.reducers.bytes.per.reducer=<number>

In order to limit the maximum number of reducers:

set hive.exec.reducers.max=<number>

In order to set a constant number of reducers:

set mapreduce.job.reduces=<number>

Starting Job = job\_1522208646737\_0003, Tracking URL =

http://localhost:8088/proxy/application\_1522208646737\_0003/

Kill Command = /usr/local/hadoop/bin/hadoop job -kill job\_1522208646737\_0003

Hadoop job information for Stage-1: number of mappers: 1; number of reducers: 1

2018-03-28 13:00:20,256 Stage-1 map = 0%, reduce = 0%

2018-03-28 13:00:28,747 Stage-1 map = 100%, reduce = 0%, Cumulative CPU 3.68 sec

2018-03-28 13:00:35,101 Stage-1 map = 100%, reduce = 100%, Cumulative CPU 6.26 sec

MapReduce Total cumulative CPU time: 6 seconds 260 msec

Ended Job = job\_1522208646737\_0003

MapReduce Jobs Launched:

Stage-Stage-1: Map: 1 Reduce: 1 Cumulative CPU: 6.26 sec HDFS Read: 9095 HDFS Write: 102 SUCCESS

Total MapReduce CPU Time Spent: 6 seconds 260 msec

OK

**35**

Time taken: 31.866 seconds, Fetched: 1 row(s)

hive>

## **d) Create index on Flight information Table**

**#create index on hbase\_flight\_new**

```
CREATE INDEX hbasefltnew_index
ON TABLE hbase_flight_new (delay)
AS 'org.apache.hadoop.hive.ql.index.compact.CompactIndexHandler'
WITH DEFERRED REBUILD;
```

SHOW INDEX ON hbase\_flight\_new;

**#create index on table hbase\_flight\_new**

```
hive> CREATE INDEX hbasefltnew_index
> ON TABLE hbase_flight_new (delay)
> AS 'org.apache.hadoop.hive.ql.index.compact.CompactIndexHandler'
> WITH DEFERRED REBUILD;
```

OK

Time taken: 0.74 seconds

### **#show index on table hbase\_flight\_new**

```
hive> SHOW INDEX ON hbase_flight_new;
```

OK

```
hbaseflightnew_index  hbase_flight_new    delay
                      default__hbase_flight_new_hbaseflightnew_index__ compact
```

Time taken: 0.104 seconds, Fetched: 1 row(s)

### **#join two tables in Hive**

#### **#create table B for join**

```
hive> create table empinfo(empno int, empgrade string) row format delimited fields terminated by
',' stored as textfile;
```

#### **#Load Data into table**

```
hive> load data local inpath '/home/hduser/Desktop/empinfo.txt' into table empinfo;
```

Loading data to table default.empinfo

OK

Time taken: 0.552 seconds

#### **#insert data into the table**

```
hive> load data local inpath '/home/hduser/Desktop/empinfo.txt' into table empinfo;
```

#### **# Table A empdbnew**

```
hive> select * from empdbnew;
```

OK

```
1    deepali120000
2    mahesh      30000
3    mangesh     25000
4    ram         39000
5    brijesh     40000
6    john        300000
```

Time taken: 0.258 seconds, Fetched: 6 row(s)

#### **# Table B empinfo**

```
hive> select * from empinfo;
```

OK

```
1    A
2    B
3    B
4    B
5    B
6    A
```

Time taken: 0.207 seconds, Fetched: 6 row(s)

#### **#Join two tables(empdbnew with empinfo on empno)**

```
hive> SELECT eno, ename, empno, empgrade FROM empdbnew JOIN empinfo ON eno = empno;
```

#### **#Join==> Result**

```

hive> SELECT eno, ename, empno, empgrade
> FROM empdbnew JOIN empinfo ON eno = empno;
WARNING: Hive-on-MR is deprecated in Hive 2 and may not be available in the future versions.
Consider using a different execution engine (i.e. spark, tez) or using Hive 1.X releases.
Query ID = hduser_20180328153258_bc345f46-a1f1-4589-ac5e-4c463834731a
Total jobs = 1
Launching Job 1 out of 1
Number of reduce tasks not specified. Estimated from input data size: 1
In order to change the average load for a reducer (in bytes):
  set hive.exec.reducers.bytes.per.reducer=<number>
In order to limit the maximum number of reducers:
  set hive.exec.reducers.max=<number>
In order to set a constant number of reducers:
  set mapreduce.job.reduces=<number>
Starting Job = job_1522208646737_0005, Tracking URL =
http://localhost:8088/proxy/application_1522208646737_0005/
Kill Command = /usr/local/hadoop/bin/hadoop job -kill job_1522208646737_0005
Hadoop job information for Stage-1: number of mappers: 2; number of reducers: 1
2018-03-28 15:33:09,615 Stage-1 map = 0%, reduce = 0%
2018-03-28 15:33:18,231 Stage-1 map = 100%, reduce = 0%, Cumulative CPU 8.17 sec
2018-03-28 15:33:24,476 Stage-1 map = 100%, reduce = 100%, Cumulative CPU 10.61 sec
MapReduce Total cumulative CPU time: 10 seconds 610 msec
Ended Job = job_1522208646737_0005
MapReduce Jobs Launched:
Stage-Stage-1: Map: 2 Reduce: 1 Cumulative CPU: 10.61 sec HDFS Read: 15336 HDFS Write:
235 SUCCESS
Total MapReduce CPU Time Spent: 10 seconds 610 msec
OK
1      deepali 1      A
2      mahesh  2      B
3      mangesh 3      B
4      ram    4      B
5      brijesh 5      B
6      john   6      A
Time taken: 26.915 seconds, Fetched: 6 row(s)

```